Fall 2022 STAT 240 In-Person (part 1) Midterm

1st Letter of Family Name | Family Name | Given Name | Student ID

Instructor (Circle) **Bret Larget** **Bi Cheng Wu**

Lecture Time (Circle) **8:50 - 9:40** **9:55 - 10:45** **1:20 - 2:10**

Discussion (Circle)

TA	time 1	time 2	time 3
Ajinkya Kokandakur	M 2:25 pm	T 7:45 am	T 4:35 pm
Cameron Jones	M 1:20 pm	T 8:50 am	T 9:55 am
Hailey Louw	T 7:45 am	T 8:50 am	T 9:55 am
Dibyendu Sah	T 8:50 am	T 9:55 am	T 11:00 am
Congwei Yang	T 8:50 am	T 9:55 am	T 11:00 am
Kessys Lorranya Peralta De Oliveira	M 2:25 pm	M 3:30 pm	
Ryan Yee	M 2:25 pm	M 4:35pm	
Margaret Turner	M 3:30 pm	M 4:35 pm	

Instructions:

1. You may use one regular sheet of paper with any handwritten notes you like. You may use both sides of the paper.
2. You may not consult other resources, your phone, a computer, information online, nor your neighbor's exam.
3. Do all of your work in the space provided. Use the backs of pages if necessary, indicating clearly that you have done so (so the grader can easily find your complete answer).

Sections

- Name, Lecture, and Discussion (2 points)
- Multiple Choice (24 points, eight questions worth 3 points each).
- Short Answer (12 points, three questions worth 4 points each).
- Data Analysis (12 points, three questions worth 4 points each).

Scoring

Problems	Possible	Score
name	2	
1 - 8	24	
9 - 11	12	
12 - 15	12	
TOTAL	50	

Name:

Grading the in-class portion of the exam means passing nearly 500 paper exams among multiple graders. We needed to organize the papers alphabetically by student within lecture section to enter grades and then sort by discussion section to return the exams. If you made our job easier by circling the requested information, you received two points. If you did not and we had to search through the course rosters to find your lecture and/or discussion section you lost one point for the lecture and one point for the discussion.

Multiple Choice (24 points). Each problem is worth 3 points.

Circle the correct answer or answers as indicated.

Problem 1. You have a dataset with a numerical column Y and a categorical (i.e. character) column X. Suppose you wanted to visualize just the Y column to see the distribution of its values in your sample. Which of the following plots would be reasonable and useful to make? **Circle all correct responses**.

(a) `geom_line()`

(b) `geom_bar()`

(c) **`geom_histogram()`**

(d) **`geom_density()`**

(e) `geom_col()`

(f) `geom_point()`

(g) `geom_smooth()`

You are summarizing a single numerical variable. (a), (f), and (g) summarize two quantitative variables. (b) and (e) summarize a single categorical variable. (c) and (d) fit the bill.

To get full credit, you needed to answer correctly. Partial credit was given if you had some answers right and others wrong.

Problem 2. Which of the following produces the proportion of values in a column X that are positive?

(a) `sum(X > 0)`

(b) **`mean(X > 0)`**

(c) `count(X > 0)`

(d) `length(X > 0)`

(a) is a sum, not a proportion; (c) would be a syntax error; (d) just returns the number of values in X whether true or false.

Problem 3. Circle **all** valid names for an R object.

(a) **`.bucky.badger.`** (b) **`bucky_badger`** (c) `bucky-badger` (d) `bucky+badger`

(a) is valid, but not best practice a the variable is hidden; (c) and (d) are invalid as they include characters which are not digits, letters, `.`, or `_`. 3 points to (a) and (b), 2 points for (a) or (b) only, 1 point for other partial credit.

Problem 4. Which command keeps all rows of the official Madison weather data set `official` where the maximum temperature `tmax` is **not missing** and the temperature is between 32 and 60 degrees Fahrenheit (include these two boundary values)?

(a) `official %>% filter( tmax != "NA" & between(tmax, 32, 60) )`
(b) `official %>% filter( tmax != "NA" && (tmax >= 32 | tmax <= 60) )`
(c) **`official %>% filter( !is.na(tmax) & tmax > 31 & tmax < 61 )`**
(d) `official %>% filter( !is.na(tmax) & (32 <= tmax <= 72) )`

Use `is.na()` to check if missing, (d) is a syntax error. The tmax variable values were all rounded to an integer (which probably should have been added to the question statement).

Problem 5. You run `df %>% group_by(type) %>% summarise(n=n(),sum=sum(x),sd=sd(x))` and get the following output:

type	n	sum	sd
A	10	81	4.6
B	7	50	5.8
C	3	29	2.7

When you run `df %>% summarise(m=mean(x))`, the final number you get is closest to which number?

(a) 5 (b) 6 (c) 7 **(d) 8** (e) 9

```
(81 + 50 + 29)
```

```
## [1] 160
```

```
(10 + 7 + 3)
```

```
## [1] 20
```

```
160 / 20
```

```
## [1] 8
```

Problem 6. Which lubridate function converts the string "10-2022-14" into the date October 14, 2022?

(a) `dmy()` (b) `dym()` **(c) `myd()`** (d) `mdy()` (e) `ymd()`

month, year, day order

Problem 7. A data set `grocery_items` has variables named `item`, `type`, and `price`. A data set named `grocery_list` has variables named `item` and `n`. The values in the columns `item` match if the same item is part of both data sets. Some items in `grocery_items` may not by in `grocery_list` and some items in `grocery_list` may not by in `grocery_items`. Which description matches the contents of `df` after executing the following code? No items are repeated within either data set. **Circle the correct response.**

```
df = grocery_list %>%
  left_join(grocery_items, by = "item")
```

(a) A data frame with one row for each item in both data sets and columns `item` and `n` only.
(b) A data frame with one row for each item in both data sets and columns `item`, `n`, `type` and `price`.
(c) **A data frame with one row for each item in `grocery_list` and columns `item`, `n`, `type` and `price`.**
(d) A data frame with one row for each item in either data set, columns `item`, `n`, `type` and `price`, and the value `NA`in columns `n`, `type`, and `price` in rows where this information was missing in the corresponding data set.

Add columns `type` and `price` from `grocery_items` to `grocery_list`, keeping all rows from `grocery list`.

Problem 8

A data set `obesity` has no missing data, 500 rows, and columns `zip`, `age_5_17`, `age_18_34`, `age_35-54`, `age_55-74`, and `age_75+`. The values in `zip` are zip codes and the values in the other columns contain counts of residents of these zip codes in the corresponding age ranges. What will be the dimensions of the data set created with the following code?

```
df = obesity %>%
  pivot_longer(starts_with("age"), names_to = "age", values_to = "n")
```

Circle all correct responses.

(a) 500 rows and 5 columns
(b) 500 rows and 6 columns
(c) **2500 rows and 3 columns**
(d) 2500 rows and 6 columns

The values in the five columns beginning with "age" are reshaped into two columns with five times as many rows, one column for the age range (`age_5_17` and so on) and one for the actual numerical values. Zip codes are repeated as needed in a single column. The number of rows is increased by a factor of 5 and reducing 5 columns to 2 reduces the total number of columns from 6 to 3.

Short Answer (12 points). Each problem is worth 4 points

Problems 9-11 are based on this small data set `df` which has numerical variables `a` and `b` and a categorical variable named `group`.

```
##   a  b group
## 1 1  3     X
## 2 2 -2     Z
## 3 3  0     Y
## 4 4  4     Z
## 5 5 -1     Z
## 6 6  5     X
```

Problem 9. Write the result of the following code.

```
df %>%
  mutate(c = a + b) %>%
  filter(group == "X") %>%
  arrange(desc(c))
```

```
df %>%
  mutate(c = a + b) %>%
  filter(group == "X") %>%
  arrange(desc(c))
```

Solution

```
## # A tibble: 2 x 4
##       a     b group     c
##   <dbl> <dbl> <chr> <dbl>
## 1     6     5 X        11
## 2     1     3 X         4
```

Problem 10 Write the result of the following code.

```
df %>%
  group_by(group) %>%
  summarize(v = min(b)) %>%
  arrange(group)
```


```
df %>%
  group_by(group) %>%
  summarize(v = min(b)) %>%
  arrange(group)
```


Solution

```
## # A tibble: 3 x 2
##   group     v
##   <chr> <dbl>
## 1 X         3
## 2 Y         0
## 3 Z        -2
```

Problem 11 Write the result of the following code.

```
df %>%
  filter(b > 0) %>%
  select(-b) %>%
  group_by(group) %>%
  mutate(v = sum(a))
```


```
df %>%
  filter(b > 0) %>%
  select(-b) %>%
  group_by(group) %>%
  mutate(v = sum(a))
```


Solution

```
## # A tibble: 3 x 3
## # Groups:   group [2]
##       a group     v
##   <dbl> <chr> <dbl>
## 1     1 X         7
## 2     4 Z         4
## 3     6 X         7
```

Data Analysis (12 points). Three problems worth 4 points each.

Each problem asks you to interpret the output from the following data analysis of the official historical Madison weather data set. The variable `prcp` measures daily precipitation in inches.

Read the questions before attempting to read the code. Only read what is needed to answer the questions.

```
## Summaries of precipitation data
prcp_1 = official %>%
  select(date, prcp) %>%
  drop_na() %>%
  mutate(year = year(date),
         month = month(date, label = TRUE),
         day = day(date),
         wday = wday(date, label = TRUE)) %>%
  group_by(year, month) %>%
  summarize(n = n(),
            v1 = max(prcp),
            v2 = sum(prcp > 0),
            v3 = sum(prcp)) %>%
  ungroup()

prcp_2 = prcp_1 %>%
  group_by(month) %>%
  summarize(w1 = min(v1),
            w2 = mean(100*v2/n),
            w3 = mean(v3))

prcp_2
```

```
## # A tibble: 12 x 4
##    month    w1    w2    w3
##    <ord> <dbl> <dbl> <dbl>
##  1 Jan    0.04  30.6  1.40
##  2 Feb    0.02  28.5  1.34
##  3 Mar    0.08  31.8  2.10
##  4 Apr    0     36.2  2.86
##  5 May    0.19  37.8  3.56
##  6 Jun    0.1   35.9  4.09
##  7 Jul    0.13  30.1  3.79
##  8 Aug    0.21  29.3  3.56
##  9 Sep    0.06  30.7  3.44
## 10 Oct    0     28.3  2.36
## 11 Nov    0.01  29.2  2.01
## 12 Dec    0.04  30.0  1.61
```

Problems 12-14 ask you what values in the summary table represent. Example solutions include the following:

- The average daily precipitation.
- The smallest total monthly precipitation in a single year.
- The average number of days per month with positive precipitation totals.

Problem 12 What does the value `w1` in the summary table represent for each month?

> `w1` is the smallest of the maximum daily precipitation totals in inches for each month. One point for each correct usage of minimum, maximum, daily, and the units inches in your response.

Problem 13 What does the value `w2` in the summary table represent for each month?

> `w2` is the percentage (or mean percentage) of days with positive precipitation for each month. Points for percentage (not precipitation) of days. Multiple points off if the description is of amounts of precipitation or for a summary of qunatities of precipitation only for days with precipitation. Percentage of precipitation lost points — what is the part and what is the whole?

Problem 14. What does the value `w3` in the summary table represent for each month?

`w3` is the mean total monthly precipitation in inches for each month. Points for each correct usage of mean, total monthly, and inches.